

Spotify Popularity Classification Analysis:

kite-luva

2025-05-24

Contents

```
# Global options for code chunks
knitr::opts_chunk$set(echo = TRUE,
                      eval = TRUE,
                      message = FALSE,
                      warning = FALSE,
                      fig.align = 'center',
                      out.width = '80%',
                      fig_caption = TRUE
                      )
```

```
#####
### CODE FOR AUC & ROC ###
#####
```

```
# Load necessary libraries
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr    1.5.1
## v ggplot2    3.5.1      v tibble     3.2.1
## v lubridate  1.9.4      v tidyr      1.3.1
## v purrr      1.0.2
```

```
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(caret)
```

```
## Loading required package: lattice
##
## Attaching package: 'caret'
##
## The following object is masked from 'package:purrr':
##
##     lift
```

```
library(ranger)
library(lubridate)
library(pROC)
```

```
## Type 'citation("pROC")' for a citation.
##
## Attaching package: 'pROC'
##
## The following objects are masked from 'package:stats':
##
##     cov, smooth, var
```

```
library(readr)
library(ggplot2)
```

```
# -----
# 1. Load & Clean the Data
# -----

# Load your dataset (adjust the file path accordingly)
spotify_charts_2024 <- read_csv("~/school docs/universal_top_spotify_songs.new.csv")
```

```
## Rows: 1750032 Columns: 25
## -- Column specification -----
## Delimiter: ","
## chr   (5): spotify_id, name, artists, country, album_name
## dbl  (17): daily_rank, daily_movement, weekly_movement, popularity, duration...
## lgl   (1): is_explicit
## date  (2): snapshot_date, album_release_date
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
# Convert date columns and calculate difference in days
spotify_charts_2024 <- spotify_charts_2024 %>%
  mutate(snapshot_date = ymd(snapshot_date),
         album_release_date = ymd(album_release_date),
         days_out = as.numeric(snapshot_date - album_release_date))

# Remove duplicates based on the spotify_id column while retaining all columns
spotify_charts_2024 <- spotify_charts_2024 %>%
  distinct(spotify_id, .keep_all = TRUE)

# Remove unneeded columns
spotify_charts_2024 <- spotify_charts_2024 %>%
  select(-country, -snapshot_date, -name, -artists, -album_name, -album_release_date, -spotify_id)

# Convert 'is_explicit' (boolean) to integer
spotify_charts_2024$is_explicit <- as.integer(spotify_charts_2024$is_explicit)

# Handle missing values in numeric columns only
numeric_cols <- sapply(spotify_charts_2024, is.numeric)
```

```

spotify_charts_2024[numeric_cols] <- lapply(spotify_charts_2024[numeric_cols],
                                             function(x) ifelse(is.na(x), median(x, na.rm = TRUE), x))

# Standardize 'duration_ms' to minutes, then remove the original column
spotify_charts_2024 <- spotify_charts_2024 %>%
  mutate(duration_min = duration_ms / 60000) %>%
  select(-duration_ms)

# -----
# 2. Prepare Data for Classification
# -----

#remove popularity 0
spotify_charts_2024 <- spotify_charts_2024 %>%
  filter(popularity != 0)

# Convert 'popularity' into a binary factor.
# This assigns popularity into two levels: "Low" and "High."
# 'make.names' ensures the levels are valid R variable names.
spotify_charts_2024 <- spotify_charts_2024 %>%
  mutate(popularity = ifelse(popularity >= 50, "High", "Low")) %>%
  mutate(popularity = make.names(popularity))

# Define feature columns (adjust these names if needed)
feature_columns <- c("daily_rank", "duration_min", "daily_movement", "weekly_movement",
                    "days_out", "is_explicit", "mode", "danceability", "energy", "loudness",
                    "speechiness", "acousticness", "instrumentalness", "time_signature",
                    "liveness", "valence", "key", "tempo")

View(spotify_charts_2024)

# Create a dataset with predictors and the target variable
class_data <- spotify_charts_2024 %>%
  select(all_of(feature_columns), popularity)

# Split the dataset into training (80%) and testing sets
set.seed(50)
trainIndex <- createDataPartition(class_data$popularity, p = 0.8, list = FALSE)
train_data <- class_data[trainIndex, ]
test_data <- class_data[-trainIndex, ]

head(train_data)

```

```

## # A tibble: 6 x 19
##   daily_rank duration_min daily_movement weekly_movement days_out is_explicit
##   <dbl>         <dbl>         <dbl>         <dbl>         <dbl>         <int>
## 1         1         4.19             0             1          191             0
## 2         3         2.96            -1             0           94             0
## 3         4         4.57            -1            -3          295             1
## 4         5         3.51             0             1          282             0
## 5         6         3.95             0             1           49             1
## 6         8         6.13             1             1           49             1
## # i 13 more variables: mode <dbl>, danceability <dbl>, energy <dbl>,
## #   loudness <dbl>, speechiness <dbl>, acousticness <dbl>,
## #   instrumentalness <dbl>, time_signature <dbl>, liveness <dbl>,

```

```
## #   valence <dbl>, key <dbl>, tempo <dbl>, popularity <chr>
```

```
# -----  
# 3. Train the Random Forest Classifier  
# -----  
  
# The 'twoClassSummary' along with 'metric = "ROC"' will use the ROC AUC for tuning.  
  
rf_model_class <- train(popularity ~ .,  
  data = train_data,  
  method = "ranger",  
  trControl = trainControl(method = "cv",  
    number = 5,  
    classProbs = TRUE,  
    summaryFunction = twoClassSummary),  
  tuneGrid = expand.grid(mtry = c(5, 7, 9, 11),  
    min.node.size = c(1, 3, 5),  
    splitrule = "gini"),  
  num.trees = 200,  
  metric = "ROC")  
print(rf_model_class)
```

```
## Random Forest  
##  
## 16117 samples  
##    18 predictor  
##    2 classes: 'High', 'Low'  
##  
## No pre-processing  
## Resampling: Cross-Validated (5 fold)  
## Summary of sample sizes: 12894, 12893, 12894, 12893, 12894  
## Resampling results across tuning parameters:  
##  
##   mtry  min.node.size  ROC      Sens      Spec  
##    5     1           0.8567575 0.9327640 0.5642857  
##    5     3           0.8567663 0.9324161 0.5642857  
##    5     5           0.8566084 0.9325032 0.5616883  
##    7     1           0.8548563 0.9311115 0.5634199  
##    7     3           0.8554429 0.9318942 0.5632035  
##    7     5           0.8562770 0.9318074 0.5658009  
##    9     1           0.8546391 0.9318074 0.5651515  
##    9     3           0.8544226 0.9309376 0.5675325  
##    9     5           0.8551210 0.9319814 0.5599567  
##   11     1           0.8532295 0.9324162 0.5634199  
##   11     3           0.8533833 0.9309376 0.5662338  
##   11     5           0.8540395 0.9307637 0.5673160  
##  
## Tuning parameter 'splitrule' was held constant at a value of gini  
## ROC was used to select the optimal model using the largest value.  
## The final values used for the model were mtry = 5, splitrule = gini  
## and min.node.size = 3.
```

```

# -----
# 4. Compute and Plot AUC & ROC Curve
# -----

# Generate predicted probabilities on the test set.
# We request probabilities (type = "prob") for both "Low" and "High" classes.
rf_pred_probs <- predict(rf_model_class, newdata = test_data, type = "prob")
head(rf_pred_probs)

##      High  Low
## 1 0.88 0.12
## 2 0.93 0.07
## 3 0.89 0.11
## 4 0.88 0.12
## 5 0.91 0.09
## 6 0.80 0.20

prob_values <- rf_pred_probs$High
# Compute the ROC curve.
# Here, we consider the probability for the "High" class as the predictor.
roc_obj <- roc(response = test_data$popularity, predictor = rf_pred_probs[, "High"])

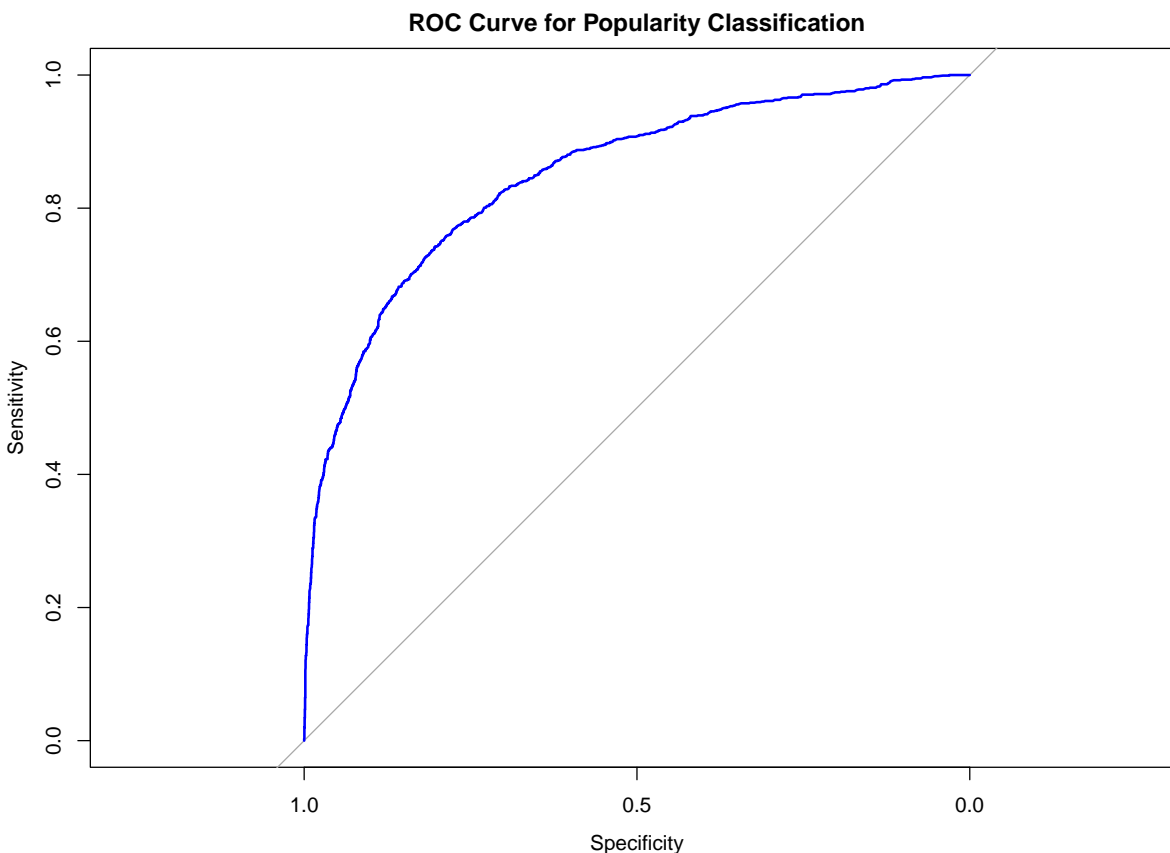
## Setting levels: control = High, case = Low
## Setting direction: controls > cases

# Calculate the AUC and print it.
auc_value <- auc(roc_obj)
cat("AUC:", auc_value, "\n")

## AUC: 0.8484169

# Plot the ROC curve.
plot(roc_obj, col = "blue", main = "ROC Curve for Popularity Classification")

```




```
#####
#####
#####
#Create an empty list to store ROC curves
roc_list <- list()
# Loop through each tuning parameter combination
for(i in 1:nrow(rf_model_class$results)){
  # extract parameters
  mtry_val <- rf_model_class$results$mtry[i]
  node_size_val <- rf_model_class$results$min.node.size[i]
  # make predictions
  predictions <- predict(rf_model_class, newdata = test_data, type= "prob")
  # Compute ROC curve
  roc_curve <- roc(test_data$popularity, predictions[, "High"])
  # Store the ROC curve with parameter labels
  roc_list[[paste("mtry=", mtry_val, " node.size=", node_size_val)]] <- roc_curve
}
```

```
## Setting levels: control = High, case = Low
## Setting direction: controls > cases
## Setting levels: control = High, case = Low
## Setting direction: controls > cases
## Setting levels: control = High, case = Low
## Setting direction: controls > cases
## Setting levels: control = High, case = Low
```

```
## Setting direction: controls > cases
## Setting levels: control = High, case = Low
## Setting direction: controls > cases
## Setting levels: control = High, case = Low
## Setting direction: controls > cases
## Setting levels: control = High, case = Low
## Setting direction: controls > cases
## Setting levels: control = High, case = Low
## Setting direction: controls > cases
## Setting levels: control = High, case = Low
## Setting direction: controls > cases
## Setting levels: control = High, case = Low
## Setting direction: controls > cases
## Setting levels: control = High, case = Low
## Setting direction: controls > cases
## Setting levels: control = High, case = Low
## Setting direction: controls > cases
```

```
###
roc_data <- do.call(rbind, lapply(names(roc_list), function(label) {
  data.frame(
    Specificity = roc_list[[label]]$specificities,
    Sensitivity = roc_list[[label]]$sensitivities,
    Combination = label
  )
}))
# Plot ROC curves with facets for each combination
ggplot(roc_data, aes(x = Specificity, y = Sensitivity, color = Combination)) +
  geom_line() +
  labs(
    title = "ROC Curves for Different mtry and min.node.size Combinations",
    x = "1-Specificity", y = "Sensitivity"
  ) +
  theme_minimal() +
  facet_wrap(~Combination) +
  scale_x_reverse()
```

ROC Curves for Different mtry and min.node.size Combinations



```
###-----

#####

# Define thresholds
upper_threshold_very_high <- 0.75
upper_threshold_high <- 0.5
lower_threshold_very_low <- 0.25

# Function to classify songs based on probability thresholds
classify_popularity <- function(prob_value) {
  if (prob_value >= upper_threshold_very_high) {
    return("Very High")
  } else if (prob_value >= upper_threshold_high) {
    return("High")
  } else if (prob_value < lower_threshold_very_low) {
    return("Very Low")
  } else if (prob_value < upper_threshold_high) {
    return("Low")
  } else {
    return("Uncertain")
  }
}

# Apply function to probability values
test_data$Predicted_Popularity <- sapply(rf_pred_probs$High, classify_popularity)
```



```

# Combine actual popularity and predicted probabilities
results_df <- data.frame(
  Actual = test_data$popularity,
  High_Probability = rf_pred_probs$High,
  Low_Probability = rf_pred_probs$Low,
  Prediction_probability = test_data$Predicted_Popularity
)

# View the first few rows
print(head(results_df))

```

```

##   Actual High_Probability Low_Probability Prediction_probability
## 1   High           0.88           0.12           Very High
## 2   High           0.93           0.07           Very High
## 3   High           0.89           0.11           Very High
## 4   High           0.88           0.12           Very High
## 5   High           0.91           0.09           Very High
## 6   High           0.80           0.20           Very High

```

```

###=====++++

```